

Sample PDF for RAG Pipeline

This is a dummy PDF file used to verify the end-to-end PDF loading, text splitting, embedding, and retrieval pipeline.

Section 1: Introduction

Retrieval Augmented Generation (RAG) combines a retrieval step over a vector store with a generative LLM. PDFs are a common source format.

Section 2: Attention Mechanism

An attention function maps a query and key-value pairs to an output. The output is a weighted sum of values where weights come from a compatibility function between the query and the corresponding key.

Section 3: Embeddings

Sentence embeddings convert text into dense vectors that capture semantic similarity. Models such as all-MiniLM-L6-v2 produce 384-d vectors suitable for nearest-neighbor retrieval.

Page 2: This page exists so the splitter has more than one page to chunk.
Multi-page handling is part of the verification.